

Quiz 8

Friday, May 17, 2024 7:48 AM

1 of 10

What is the main difference between **passive reinforcement learning** and evaluating a policy π by **solving an MDP** (Markov Decision Process) model?

- A** There is no difference because passive reinforcement learning is also to evaluate a policy π .

B Passive reinforcement learning can find the optimal policy.

C In passive reinforcement learning, the transition model and the reward function are unknown.

D Passive reinforcement learning can learn the action-value function $Q(s, a)$ while Policy Evaluation with MDP can learn the state-value function $V(s)$.

SUBMIT ANSWER

2 of 10

```
Initialize:
   $V(s) \in \mathbb{R}$ , arbitrarily, for all  $s \in \mathcal{S}$ 
   $Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$ 
Loop forever (for each episode):
  Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$ 
   $G \leftarrow 0$ 
  Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :
     $G \leftarrow \gamma G + R_{t+1}$ 
    Unless  $S_t$  appears in  $S_0, S_1, \dots, S_{t-1}$ :
      Append  $G$  to  $Returns(S_t)$ 
     $V(S_t) \leftarrow \text{average}(Returns(S_t))$ 
```

Zoom

What does the following algorithm do?

- A** Estimating a state-value function for a policy π .

B Estimating an action-value function for a policy π .

C Estimating the optimal state-value function.

D Estimating the optimal action-value function.

SUBMIT ANSWER

```
Initialize:
   $V(s) \in \mathbb{R}$ , arbitrarily, for all  $s \in \mathcal{S}$ 
   $Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$ 
Loop forever (for each episode):
  Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$ 
   $G \leftarrow 0$ 
  Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :
     $G \leftarrow \gamma G + R_{t+1}$ 
    Unless  $S_t$  appears in  $S_0, S_1, \dots, S_{t-1}$ :
      Append  $G$  to  $Returns(S_t)$ 
     $V(S_t) \leftarrow \text{average}(Returns(S_t))$ 
```

3 of 10

Input Policy π

Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1	Episode 2
B, right, C, -1 C, up, A, -1 A, exit, x, +8	B, right, C, -1 C, up, A, -1 A, exit, x, +8
Episode 3	Episode 4
E, up, C, -1 C, up, A, -1 A, exit, x, +8	E, up, C, -1 C, up, D, -1 D, exit, x, -10

Zoom

Given the following environment with five states A, B, C, D, and E where A and D are two terminal states.

Let's consider the following policy π : $\pi(A)=\text{exit}$, $\pi(B)=\text{right}$, $\pi(C)=\text{up}$, $\pi(D)=\text{exit}$, $\pi(E)=\text{up}$.

We use the **Monte Carlo Direct Evaluation** method to evaluate the policy π . We obtain **four episodes** as in the below picture.

Assuming the discount factor $\gamma=1$, **compute the expected utility of the five states $V(A)=?$, $V(B)=?$, $V(C)=?$, $V(D)=?$, $V(E)=?$**

- A** $V(A) = 8, V(B) = 6, V(C)=2, V(D) = -10, V(E) = -6$
- B** $V(A) = 8, V(B) = 6, V(C)=2.5, V(D) = -10, V(E) = -3$
- C** $V(A) = 6, V(B) = 3, V(C)=2.5, V(D) = -2.5, V(E) = -1.5$
- D** $V(A) = 8, V(B) = 6, V(C)=2, V(D) = -10, V(E) = -6$

SUBMIT ANSWER

4 of 10

Algorithm parameter: step size $\alpha \in (0,1]$
 Initialize $V(s)$, for all $s \in S^+$, arbitrarily except that $V(\text{terminal}) = 0$
 Loop for each episode:
 Initialize S
 Loop for each step of episode:
 $A \leftarrow$ action given by π for S
 Take action A , observe R, S'
 $V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$
 $S \leftarrow S'$
 until S is terminal

Zoom

What is the following algorithm?

- A** A temporal difference learning algorithm to estimate the state-value function for a policy π .
- B** A Monte Carlo learning algorithm to estimate the state-value function for a policy π .
- C** SARSA algorithm for estimating the optimal state-value function.
- D** SARSA algorithm for estimating the optimal action-value function.


SUBMIT ANSWER

```

Algorithm parameter: step size  $\alpha \in (0, 1]$ 
Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$ 
Loop for each episode:
  Initialize  $S$ 
  Loop for each step of episode:
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe  $R, S'$ 
     $V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal

```

5 of 10



Given the following environment with five states A, B, C, D, and E where A and D are two terminal states.

Let's consider the following

policy π : $\pi(A)=\text{exit}$, $\pi(B)=\text{right}$, $\pi(C)=\text{right}$, $\pi(D)=\text{exit}$, $\pi(E)=\text{up}$.

We use the **Temporal Difference Learning** method to evaluate the policy π .

We have two transitions as in the below picture. (B,right,C,-4) gives us a reward of -4. (C, right, D,-2) gives us a reward of -2.

Assuming the discount factor $\gamma=1$, learning rate $\alpha=1/2$, **compute the expected utility of the five states after the two transitions** $V(A)=?$, $V(B)=?$, $V(C)=?$, $V(D)=?$, $V(E)=?$

A

$V(A)=0$, $V(B)=-2$, $V(C)=2$, $V(D)=6$, $V(E)=0$

B

$V(A)=0$, $V(B)=-2$, $V(C)=-2$, $V(D)=3$, $V(E)=0$

C

$V(A)=-2$, $V(B)=2$, $V(C)=2$, $V(D)=6$, $V(E)=0$

D

$V(A)=0$, $V(B)=2$, $V(C)=-2$, $V(D)=6$, $V(E)=0$

6 of 10

What is the key difference between **passive** reinforcement learning and **active** reinforcement learning?

A

Passive RL is to learn the value function for a fixed policy π while active RL is to learn the optimal policy/value function.

B

The transition function $P(s'|s, a)$ is unknown in active RL.

C

The reward function $R(s, a, s')$ is unknown in active RL.

D

We can employ the Value Iteration/Policy Iteration algorithms to solve active RL.

SUBMIT ANSWER

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$
 $Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$
 Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 $G \leftarrow \gamma G + R_{t+1}$
 Unless S_t appears in S_0, S_1, \dots, S_{t-1} :
 Append G to $Returns(S_t)$
 $V(S_t) \leftarrow \text{average}(Returns(S_t))$

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$
 $Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$
 Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 $G \leftarrow \gamma G + R_{t+1}$
 Unless S_t appears in S_0, S_1, \dots, S_{t-1} :
 Append G to $Returns(S_t)$
 $V(S_t) \leftarrow \text{average}(Returns(S_t))$

Zoom

How should we change the following Reinforcement Learning algorithm so that we can estimate an optimal value function?

- ☐ A Estimating the action-value function $Q(S_t, A_t)$ instead of the state-value function $V(S_t)$.
- ☐ B Performing one-step lookahead to improve the current estimated state-value function
 $V(S_t) \leftarrow \max_a P(S_{t+1}|S_t, A_t)[R(S_t, A_t, S_{t+1}) + \gamma V(S_{t+1})]$
- ☐ C Performing policy improvement with respect to the current estimated action-value function, i.e.,
 $\pi(S_t) \leftarrow \arg \max_a Q(S_t, a)$.
- ☐ D Choose $S_0 \in \mathcal{S}$, and $A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0 .
- ☐ E A, B, and C need to be done.
- ☒ F A, C, and D need to be done.
- ☐ G B and D need to be done.
- ☐ H A, B, C, and D need to be done.

8 of 10

What is the $Q(s, a)$ update rule in Q-Learning? α is the learning rate, γ is the discount factor.

- ☐ A $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma Q(s', a'))$
- ☒ B $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$
- ☐ C $Q(s, a) \leftarrow \alpha(r + \gamma \max_{a'} Q(s', a')) - Q(s, a)$
- ☐ D $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$
- ☐ E All are correct.

SUBMIT ANSWER

9 of 10


A is the set of actions a 's we can perform at each state $s \in S$. Let $n = |A|$ be the number of actions.

An ϵ -greedy policy would choose a random action with probability ϵ .

Regarding an ϵ -greedy policy with respect to a action-value function $Q(s, a)$, in each step of **Q-Learning with ϵ -greedy exploration**, what is the probability that the action with the maximum $Q(s, a)$ value would be sampled?

- A
- B
- C
- D

SUBMIT ANSWER




We are using **Approximate Q-Learning** with linear Q-value approximation for a basic Pacman game with two features: the reciprocal distance to the nearest food dot f_{DOT} (with weight w_{DOT}) and the reciprocal distance to the nearest ghost f_{GST} (with weight w_{GST}).

Let's consider the current state s and action NORTH (going up) in the picture below. After Pacman performs the action NORTH, the blue ghost attacks Pacman, the obtained reward is -1000 in this transition, and the game ends.

Question: Using this transition, how the weights w_{DOT} and w_{GST} are updated?

Note: discount factor $\gamma = 1$, and learning rate $\alpha = 0.001$.

- A
- B
- C
- D



Finished!

Score: 9/10

Percent: 90%

OK